

```
In [1]: import pandas as pd
import numpy as np
```

```
In [2]: df = pd.read_csv('diabetes.csv')
df
```

Out[2]:

	Pregnancies	Glucose	BloodPressure	SkinThickness	Insulin	BMI	DiabetesPedigreeFunction	Age	Outcome
0	6	148	72	35	0	33.6	0.627	50	1
1	1	85	66	29	0	26.6	0.351	31	0
2	8	183	64	0	0	23.3	0.672	32	1
3	1	89	66	23	94	28.1	0.167	21	0
4	0	137	40	35	168	43.1	2.288	33	1
...
763	10	101	76	48	180	32.9	0.171	63	0
764	2	122	70	27	0	36.8	0.340	27	0
765	5	121	72	23	112	26.2	0.245	30	0
766	1	126	60	0	0	30.1	0.349	47	1
767	1	93	70	31	0	30.4	0.315	23	0

768 rows × 9 columns

In [3]: df.info()

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 768 entries, 0 to 767
Data columns (total 9 columns):
#   Column                Non-Null Count  Dtype
---  -
0   Pregnancies           768 non-null    int64
1   Glucose               768 non-null    int64
2   BloodPressure         768 non-null    int64
3   SkinThickness         768 non-null    int64
4   Insulin               768 non-null    int64
5   BMI                  768 non-null    float64
6   DiabetesPedigreeFunction 768 non-null    float64
7   Age                  768 non-null    int64
8   Outcome               768 non-null    int64
dtypes: float64(2), int64(7)
memory usage: 54.1 KB
```

In [4]: df.head()

Out[4]:

	Pregnancies	Glucose	BloodPressure	SkinThickness	Insulin	BMI	DiabetesPedigreeFunction	Age	Outcome
0	6	148	72	35	0	33.6	0.627	50	1
1	1	85	66	29	0	26.6	0.351	31	0
2	8	183	64	0	0	23.3	0.672	32	1
3	1	89	66	23	94	28.1	0.167	21	0
4	0	137	40	35	168	43.1	2.288	33	1

In [5]: df.Outcome.value_counts()

Out[5]: 0 500
1 268
Name: Outcome, dtype: int64

```
In [6]: df1 = df[df['Outcome'] == 0].iloc[:300,:]  
df2 = df[df['Outcome'] == 1]
```

```
In [7]: df3 = pd.concat([df1,df2],axis=0)
```

```
In [8]: df3.Outcome.value_counts()
```

```
Out[8]: 0    300  
       1    268  
       Name: Outcome, dtype: int64
```

```
In [9]: x,y = df.iloc[:, :-1], df.Outcome
```

```
In [10]: from sklearn.preprocessing import MinMaxScaler
```

```
In [11]: mn = MinMaxScaler()  
x = mn.fit_transform(x)  
x
```

```
Out[11]: array([[0.35294118, 0.74371859, 0.59016393, ..., 0.50074516, 0.23441503,  
                0.48333333],  
               [0.05882353, 0.42713568, 0.54098361, ..., 0.39642325, 0.11656704,  
                0.16666667],  
               [0.47058824, 0.91959799, 0.52459016, ..., 0.34724292, 0.25362938,  
                0.18333333],  
               ...,  
               [0.29411765, 0.6080402 , 0.59016393, ..., 0.390462 , 0.07130658,  
                0.15      ],  
               [0.05882353, 0.63316583, 0.49180328, ..., 0.4485842 , 0.11571307,  
                0.43333333],  
               [0.05882353, 0.46733668, 0.57377049, ..., 0.45305514, 0.10119556,  
                0.03333333]])
```

```
In [12]: from sklearn.model_selection import train_test_split
```

```
In [13]: xtrain,xtest,ytrain,ytest = train_test_split(x,y,test_size=0.2,random_state=2)
```

```
In [14]: from sklearn.ensemble import RandomForestClassifier, AdaBoostClassifier, GradientBoostingClassifier
from xgboost import XGBClassifier
```

```
In [15]: from sklearn.metrics import classification_report,f1_score
```

```
In [16]: abc = AdaBoostClassifier()
abc.fit(xtrain,ytrain)
ypred = abc.predict(xtest)

print(classification_report(ytest,ypred))
print(f1_score(ytest,ypred,average=None,labels=[0,1]))
```

	precision	recall	f1-score	support
0	0.78	0.83	0.81	109
1	0.53	0.44	0.48	45
accuracy			0.72	154
macro avg	0.66	0.64	0.65	154
weighted avg	0.71	0.72	0.71	154

```
[0.80888889 0.48192771]
```

```
In [17]: gbc = GradientBoostingClassifier()
gbc.fit(xtrain,ytrain)
ypred = gbc.predict(xtest)

print(classification_report(ytest,ypred))
print(f1_score(ytest,ypred,average=None,labels=[0,1]))
```

	precision	recall	f1-score	support
0	0.81	0.81	0.81	109
1	0.54	0.56	0.55	45
accuracy			0.73	154
macro avg	0.68	0.68	0.68	154
weighted avg	0.74	0.73	0.73	154

[0.81105991 0.54945055]

```
In [18]: xgb = XGBClassifier()
xgb.fit(xtrain,ytrain)
ypred = xgb.predict(xtest)

print(classification_report(ytest,ypred))
print(f1_score(ytest,ypred,average=None,labels=[0,1]))
```

	precision	recall	f1-score	support
0	0.82	0.79	0.80	109
1	0.53	0.58	0.55	45
accuracy			0.73	154
macro avg	0.67	0.68	0.68	154
weighted avg	0.73	0.73	0.73	154

[0.80373832 0.55319149]

```
In [19]: rf = RandomForestClassifier()
rf.fit(xtrain, ytrain)
ypred = rf.predict(xtest)

print(classification_report(ytest,ypred))
print(f1_score(ytest,ypred,average=None,labels=[0,1]))
```

	precision	recall	f1-score	support
0	0.82	0.85	0.84	109
1	0.61	0.56	0.58	45
accuracy			0.77	154
macro avg	0.72	0.70	0.71	154
weighted avg	0.76	0.77	0.76	154

[0.83783784 0.58139535]

Balancing target column

```
In [20]: df1 = df[df['Outcome'] == 0].iloc[:300,: ]
df2 = df[df['Outcome'] == 1]

df3 = pd.concat([df1,df2],axis=0)

df3.Outcome.value_counts()
```

```
Out[20]: 0    300
1    268
Name: Outcome, dtype: int64
```

```
In [21]: x,y = df3.iloc[:,-1],df3.Outcome
```

```
In [22]: mn = MinMaxScaler()  
x = mn.fit_transform(x)  
x
```

```
Out[22]: array([[0.05882353, 0.42713568, 0.54098361, ..., 0.39642325, 0.11656704,  
                0.16666667],  
               [0.05882353, 0.44723618, 0.54098361, ..., 0.41877794, 0.03800171,  
                0.          ],  
               [0.29411765, 0.58291457, 0.60655738, ..., 0.38152012, 0.05251921,  
                0.15        ],  
               ...,  
               [0.35294118, 0.95477387, 0.75409836, ..., 0.5290611 , 0.0853971 ,  
                0.75        ],  
               [0.52941176, 0.85427136, 0.60655738, ..., 0.6557377 , 0.13877028,  
                0.36666667],  
               [0.05882353, 0.63316583, 0.49180328, ..., 0.4485842 , 0.11571307,  
                0.43333333]])
```

```
In [23]: xtrain,xtest,ytrain,ytest = train_test_split(x,y,test_size=0.20,random_state=2)
```

```
In [24]: abc = AdaBoostClassifier()
abc.fit(xtrain,ytrain)
ypred = abc.predict(xtest)

print(classification_report(ytest,ypred))
print(f1_score(ytest,ypred,average=None,labels=[0,1]))
```

	precision	recall	f1-score	support
0	0.80	0.70	0.75	61
1	0.70	0.79	0.74	53
accuracy			0.75	114
macro avg	0.75	0.75	0.75	114
weighted avg	0.75	0.75	0.75	114

[0.74782609 0.74336283]

```
In [25]: gbc = GradientBoostingClassifier()
gbc.fit(xtrain,ytrain)
ypred = gbc.predict(xtest)

print(classification_report(ytest,ypred))
print(f1_score(ytest,ypred,average=None,labels=[0,1]))
```

	precision	recall	f1-score	support
0	0.77	0.67	0.72	61
1	0.67	0.77	0.72	53
accuracy			0.72	114
macro avg	0.72	0.72	0.72	114
weighted avg	0.73	0.72	0.72	114

[0.71929825 0.71929825]


```
In [26]: xgb = XGBClassifier()
xgb.fit(xtrain,ytrain)
ypred = xgb.predict(xtest)

print(classification_report(ytest,ypred))
print(f1_score(ytest,ypred,average=None,labels=[0,1]))
```

	precision	recall	f1-score	support
0	0.71	0.66	0.68	61
1	0.64	0.70	0.67	53
accuracy			0.68	114
macro avg	0.68	0.68	0.68	114
weighted avg	0.68	0.68	0.68	114

[0.68376068 0.66666667]

```
In [27]: rf = RandomForestClassifier()
rf.fit(xtrain, ytrain)
ypred = rf.predict(xtest)

print(classification_report(ytest,ypred))
print(f1_score(ytest,ypred,average=None,labels=[0,1]))
```

	precision	recall	f1-score	support
0	0.75	0.69	0.72	61
1	0.67	0.74	0.70	53
accuracy			0.71	114
macro avg	0.71	0.71	0.71	114
weighted avg	0.71	0.71	0.71	114

[0.71794872 0.7027027]

Adaboost is performing better

```
In [28]: abc = AdaBoostClassifier()
abc.fit(xtrain,ytrain)
ypred = abc.predict(xtest)

print(classification_report(ytest,ypred))
print(f1_score(ytest,ypred,average=None,labels=[0,1]))
```

	precision	recall	f1-score	support
0	0.80	0.70	0.75	61
1	0.70	0.79	0.74	53
accuracy			0.75	114
macro avg	0.75	0.75	0.75	114
weighted avg	0.75	0.75	0.75	114

[0.74782609 0.74336283]

```
In [29]: print(abc.score(xtrain, ytrain))
print(abc.score(xtest, ytest))
```

0.8480176211453745
0.7456140350877193

```
In [30]: # for i in range(5000,10000):
#         xtrain,xtest,ytrain,ytest = train_test_split(x,y,test_size=0.25,random_state=i)
#         abc = AdaBoostClassifier()
#         abc.fit(xtrain,ytrain)
#         ypred = abc.predict(xtest)
#         if abc.score(xtrain, ytrain) > 0.82 and abc.score(xtest, ytest) > 0.81:
#             print('Random State:', i)
#             print(abc.score(xtrain, ytrain))
#             print(abc.score(xtest, ytest))
#             print()
```

```
In [31]: xtrain,xtest,ytrain,ytest = train_test_split(x,y,test_size=0.25,random_state=7998)
abc = AdaBoostClassifier()
abc.fit(xtrain,ytrain)
ypred = abc.predict(xtest)

print(abc.score(xtrain, ytrain))
print(abc.score(xtest, ytest))

print(classification_report(ytest,ypred))
print(f1_score(ytest,ypred,average=None,labels=[0,1]))
```

0.823943661971831

0.823943661971831

	precision	recall	f1-score	support
0	0.83	0.84	0.83	75
1	0.82	0.81	0.81	67
accuracy			0.82	142
macro avg	0.82	0.82	0.82	142
weighted avg	0.82	0.82	0.82	142

[0.83443709 0.81203008]

In []:

In []:

In []: